



UNIVERSITY
INSTITUTE
OF LISBON

A Cross-Layer SDN Orchestration Architecture for Stateful Edge Services

Pedro Ojo Oriakhi

Master's in Telecommunications and Computer Engineering

Supervisors:

Ph.D. Professor José Moura, Assistant Professor,
Iscte - Instituto Universitário de Lisboa

Ph.D. Professor Pedro Ramos, Assistant Professor,
Iscte - Instituto Universitário de Lisboa

Month, Year

[This page is intentionally left blank.]

Department of Information Science and Technology

A Cross-Layer SDN Orchestration Architecture for Stateful Edge Services

Pedro Ojo Oriakhi

Master's in Telecommunications and Computer Engineering

Supervisors:

Ph.D. Professor José Moura, Assistant Professor,
Iscte - Instituto Universitário de Lisboa

Ph.D. Professor Pedro Ramos, Assistant Professor,
Iscte - Instituto Universitário de Lisboa

Month, Year

[This page is intentionally left blank.]

Write here your dedication.

[This page is intentionally left blank.]

Acknowledgments

Write here your acknowledgements and, if applicable, any grants or sources of funding.

[This page is intentionally left blank.]

Resumo

Palavras-Chave:

[This page is intentionally left blank.]

Abstract

Write here the abstract in English. There is a 250 word limit.

Keywords: *Keywords; separated; by; semicolons*

[This page is intentionally left blank.]

Contents

Acknowledgments	iii
Resumo	v
Abstract	vii
List of Figures	xi
List of Tables	xiii
List of Acronyms	xv
Chapter 1. Introduction	1
1.1. Context, Motivation, and Problem Statement	1
1.2. Objectives	3
1.3. Research Questions	3
1.4. Research Methodology	3
1.5. Main Contributions of this Dissertation	3
1.6. Dissertation Structure	3
Chapter 2. Background and Related Work	5
2.1. Review Methodology	5
2.2. The Unexamined Default: Three-Layer Separation	5
2.3. Auto Scaling	5
2.4. Load Balancing on SDN	5
2.5. Monitoring & Telemetry	5
2.6. Resource Orchestration on SDN	5
2.7. Summary and Research Gaps	5
Chapter 3. Architecture and Design of Proposed Solution for Edge Services	7
3.1. Design Requirements	7
3.2. Architecture	7
3.3. Distributed Elastic Allocation of Resources	7
3.3.1. Services	7
3.3.2. Data	7
3.4. Monitoring and Decision Engine	7
3.5. Control Workflow	7
Chapter 4. Implementation of Proposed Solution	9

4.1. Experimental infrastructure	9
4.2. Distributed Elastic Allocation of Resources	9
4.2.1. Services	9
4.2.2. Data	9
4.3. Monitoring and Decision Engine	9
4.4. Control Workflow	9
4.5. Implementation Validation	9
Chapter 5. Evaluation Methodology	11
5.1. Experimental Objectives	11
5.2. Evaluation Scenarios	11
5.3. Performance Metrics	11
5.4. Experimental Procedure with Results Statistical Analysis	11
Chapter 6. Experimental Results	13
6.1. Telemetry Delivery Semantics (RQ1)	13
6.2. Bottleneck-Aware Scaling Action (RQ2)	13
6.3. Readiness Propagation and Traffic Admission (RQ3)	13
6.4. Network Performance	13
6.5. Scalability Analysis	13
6.6. Discussion: Synthesis Across Interfaces	13
Chapter 7. Conclusions and Future Work	15
7.1. Conclusions	15
7.2. Research Contributions	15
7.3. Limitations	15
7.4. Future Research Directions	15
References	17

List of Figures

Figure 1.1 Global internet traffic between 2020 and 2025. Source: Statista, based on ITU data [1].

2

[This page is intentionally left blank.]

List of Tables

[This page is intentionally left blank.]

List of Acronyms

:

[This page is intentionally left blank.]

CHAPTER 1

Introduction

Studies have shown that edge web services encompassing routing, monitoring and auto-scaling can suffer from coordination delays that arise from the independent control loops through which each module operates. This accumulation of handoff delays between components degrades service quality and can reduce orchestration reliability. In this chapter the context, motivation and problem statement are explained in 1.1 with the objectives presented in 1.2. The research questions that emerged are available in 1.3. The methodology used to develop the thesis is detailed in 1.4, the main contributions in 1.5 and finally the outlining of the rest of the 7 chapters of the thesis are presented in 1.6.

1.1. Context, Motivation, and Problem Statement

In the year 2025 global internet traffic surpassed 7 exabytes (1 exabyte = 1 000 000 000 gigabytes), this corresponds to a more than double increase in the amount of internet traffic generated since 2020 across fixed and mobile networks, this is represented in Figure 1.1. [1].

This growth has been driven by many technological advancements some of which, have in turn galvanized the emergence of real-time interactive and data-intensive applications among others. Examples include web streaming and cloud gaming, which are both services where users expect sub-second latency and high resolution multimedia, highlighting the need for fast and capable data processing. These are just a few examples that demonstrate the demand of infrastructure that is able to deliver performance with high availability and reliability at global scale, in other words regardless of where the end users are located.

Cloud computing is and has been the [2] backbone that enables these applications to operate at a global scale. It offers automatic scaling, both horizontally and vertically, a pay-per-use model and among other aspects, [2] a globally distributed offering and availability that has established it as the foundational layer that supports modern applications and services.

But at the same time, a growing percentage of global internet traffic comes from mobile devices, which now represent the main gateway to internet access [3]. This ongoing shift underscores that the majority of internet access comes from the network edge, far from the cloud data centers, evidencing the importance of edge computing as an extension or alternative to cloud computing as it enables geo-distributed deployments of applications and services. As presented in [4], [5] edge computing brings the computational resources to the network edge which and therefore closer to the end-users and end-products allowing for much lower latency in provisioning services and applications to users. Beyond lower




FIGURE 1.1. Global internet traffic between 2020 and 2025. Source: Statista, based on ITU data [1].

latency, edge computing's inherent geo-distributed servers dispersed geographically and close to users relieves bandwidth pressure on the core network, since data can be processed and stored locally, at the edge, for quicker responses, only a subset of it is sent to the cloud [4], freeing up backbone available bandwidth.

Despite the advantages edge computing also brings concerns regarding aspects such as resources scarcity as per-site capex. And the management of distributed edge servers, as Satyanarayanan said in [5] "The dispersion inherent in edge computing raises the complexity of management considerably" Resource provision is important because as in [6], there is variation in user requests overtime demanding the edge infrastructure to be flexible to adapt to these fluctuations in an efficient way.

So focus on the complexity of managing edge resource/edge servers which is because of the heterogeneity in requests, varying demand and distributed nature. You can focus on the ones that matter to you

On the flip side, edge computing is not a replacement for cloud computing [4], rather a complement or extension as cloud computing can be understood as the persistent computing and storage layer to be

- 1.2. Objectives
- 1.3. Research Questions
- 1.4. Research Methodology
- 1.5. Main Contributions of this Dissertation
- 1.6. Dissertation Structure

[This page is intentionally left blank.]

CHAPTER 2

Background and Related Work

- 2.1. Review Methodology
- 2.2. The Unexamined Default: Three-Layer Separation
- 2.3. Auto Scaling
- 2.4. Load Balancing on SDN
- 2.5. Monitoring & Telemetry
- 2.6. Resource Orchestration on SDN
- 2.7. Summary and Research Gaps

[This page is intentionally left blank.]

CHAPTER 3

Architecture and Design of Proposed Solution for Edge Services

This chapter presents the design of the proposed cross layer sdn orchestration architecture for edge services. It stems from the research gaps explored in the 2 and brings to the fold the need for an architecture that allows for each mentioned component to be independently managed and tested.

3.1. Design Requirements

3.2. Architecture

3.3. Distributed Elastic Allocation of Resources

3.3.1. Services

3.3.2. Data

3.4. Monitoring and Decision Engine

3.5. Control Workflow

[This page is intentionally left blank.]

CHAPTER 4

Implementation of Proposed Solution

4.1. Experimental infrastructure

4.2. Distributed Elastic Allocation of Resources

4.2.1. Services

4.2.2. Data

4.3. Monitoring and Decision Engine

4.4. Control Workflow

4.5. Implementation Validation

[This page is intentionally left blank.]

CHAPTER 5

Evaluation Methodology

5.1. Experimental Objectives

5.2. Evaluation Scenarios

5.3. Performance Metrics

5.4. Experimental Procedure with Results Statistical Analysis

[This page is intentionally left blank.]

CHAPTER 6

Experimental Results

- 6.1. Telemetry Delivery Semantics (RQ1)
- 6.2. Bottleneck-Aware Scaling Action (RQ2)
- 6.3. Readiness Propagation and Traffic Admission (RQ3)
- 6.4. Network Performance
- 6.5. Scalability Analysis
- 6.6. Discussion: Synthesis Across Interfaces

[This page is intentionally left blank.]

CHAPTER 7

Conclusions and Future Work

7.1. Conclusions

7.2. Research Contributions

7.3. Limitations

7.4. Future Research Directions

[This page is intentionally left blank.]

References

- [1] International Telecommunication Union, “Facts and figures 2025: Internet traffic,” ITU, Oct. 2025. Accessed: Jul. 13, 2026. [Online]. Available: <https://www.itu.int/itu-d/reports/statistics/2025/10/15/ff25-internet-traffic/>
- [2] D. Gurung, S. M. Z. U. Rashid, Z. u. Abdeen, and S. Rath, “Cloud revolution: Tracing the origins and rise of cloud computing,” in *2026 IEEE 16th Annual Computing and Communication Workshop and Conference (CCWC)*, Las Vegas, NV, USA, Jan. 2026, pp. 1100–1106. DOI: 10.1109/CCWC67433.2026.11393790 [Online]. Available: <https://doi.org/10.1109/CCWC67433.2026.11393790>
- [3] T. Gaudiaut, *Mobile accounts for nearly 60 percent of web traffic*, Statista Chart of the Day, Data from StatCounter. Licensed under CC BY-ND 3.0, Mar. 2026. Accessed: Jul. 19, 2026. [Online]. Available: <https://www.statista.com/chart/35951/share-of-mobile-device-website-traffic-worldwide-since-2010/>
- [4] K. Cao, Y. Liu, G. Meng, and Q. Sun, “An overview on edge computing research,” *IEEE Access*, vol. 8, pp. 85 714–85 728, 2020. DOI: 10.1109/ACCESS.2020.2991734 [Online]. Available: <https://doi.org/10.1109/ACCESS.2020.2991734>
- [5] M. Satyanarayanan, “The emergence of edge computing,” *Computer*, vol. 50, no. 1, pp. 30–39, Jan. 2017. DOI: 10.1109/MC.2017.9 [Online]. Available: <https://doi.org/10.1109/MC.2017.9>
- [6] Q. Luo, S. Hu, C. Li, G. Li, and W. Shi, “Resource scheduling in edge computing: A survey,” *IEEE Communications Surveys & Tutorials*, vol. 23, no. 4, pp. 2131–2165, 2021. DOI: 10.1109/COMST.2021.3106401 [Online]. Available: <https://doi.org/10.1109/COMST.2021.3106401>